

Classification of musical genres

Final project for Applied Machine Learning 2020

Eliot, Mads and Sofus

Finding features

We use the GTZAN dataset [1]


- 1000 songs in 10 genres: pop, jazz, blues, classical, hiphop, metal, rock, reggae, disco, country
- Decided it was fun to find our own features and use a bunch of different approaches
- Use Timbre and Librosa python packages

train_features =

Librosa: zerocrossing, spectralrolloff, chromagramstd, chromagrammean, chroma_cens

Timbre: hardness, depth, brightness, roughness, warmth, sharpness, boominess

Additional: fourier_std, fourier_mean




[1] <https://www.kaggle.com/andradaolteanu/gtzan-dataset-music-genre-classification>

The initial approach: Gradient boosted tree and Nearest Neighbor

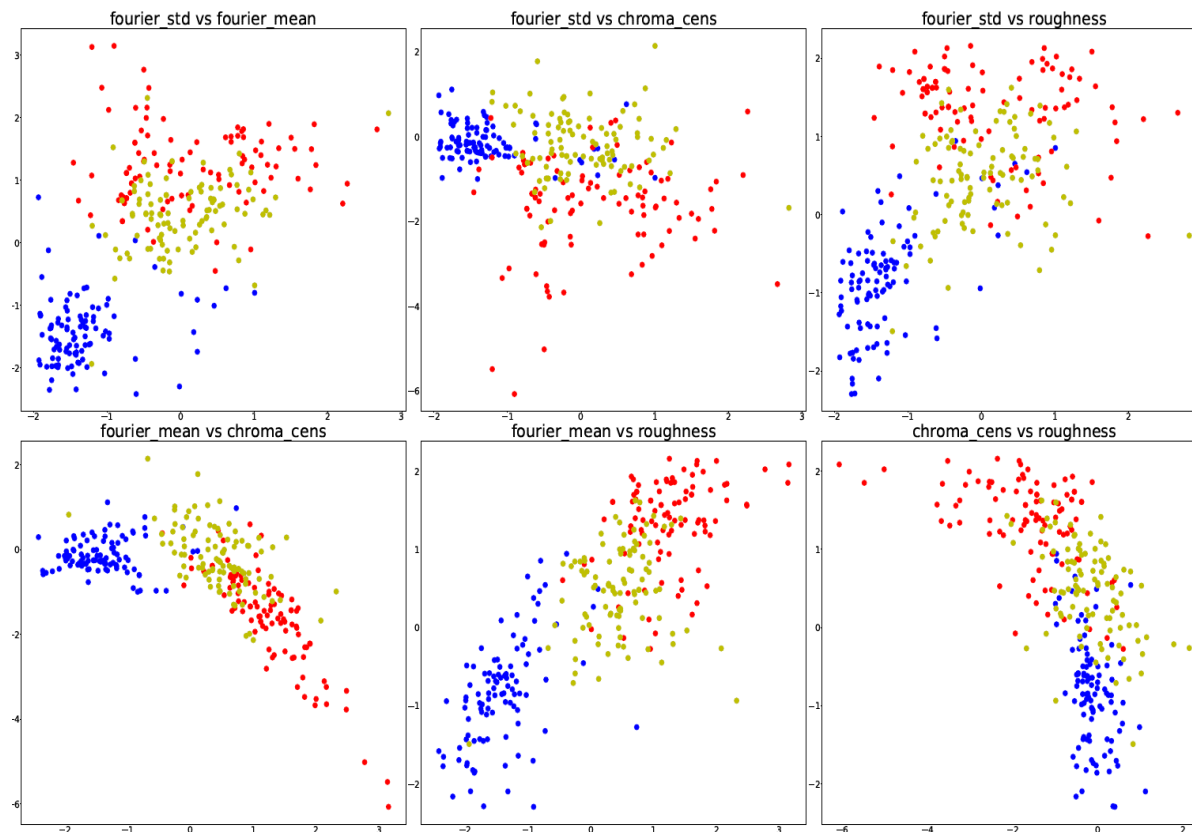
For three genres, kNeighbour gives an accuracy score of 0.99 ± 0.04 with 50 fold cross validation

- This drops to 0.55 ± 0.11 for 10 genres

For three genres, LightGBM classifier gives an accuracy score of 0.80 ± 0.03 , but holds up better for more genres.

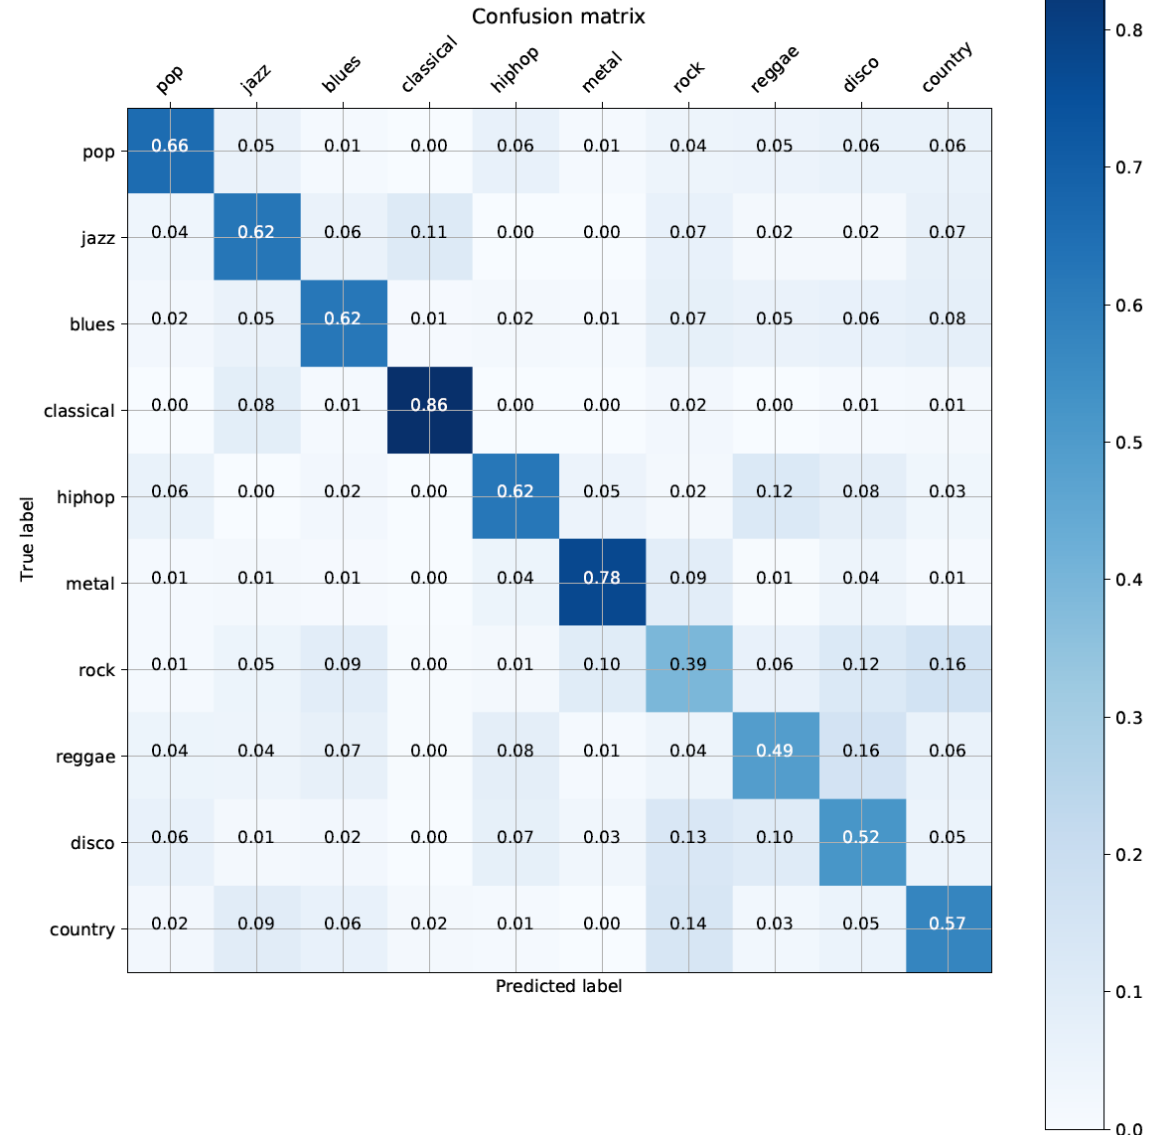


- Metal
- Disco
- Classical





Finding errors in the prediction – what genres are hardest to predict?

- Classical and metal are easiest to predict
- Rock is hardest, which makes sense intuitively – vaguely defined!
- Can we introduce variables that improves the guesses on the lowest scoring classes?



Harmonic Percussive Source Separation





- Median filtering of spectrogram
- Analyse melodic and rhythmic features separately



For details on HPSS, see: Derry FitzGerald Proc. of the 13th Int. Conference on Digital Audio Effects (DAFx-10) (2010)

Rythmic features



- **Tempo** : estimate from autocorrelation
- **onBeat** : $\frac{\int \delta(t-t_n) OE(t) dt}{\int OE(t) dt}$
- **Overlap** : $\int PSD[OE](f) \text{norm}(f; \mu_i, \sigma_i) df$
 - Early and Late
- **Synco** : Early - Late



Feature Importance

- Permutation (10 repeats) and split importance
- SHAP values


... Homemade features are competitive!




Improvement?

- LightGBM Classifier (log loss)
- StandardScaler, Data split Train:Val:Test = 70:10:20
- Hyper parameter optimization (GridSearch with 5CV)
 - $\Gamma_{\text{learn}} = 8.14 \times 10^{-4}$, max depth = 6, num leaves = 9, $n_{\text{estimators}} = 30000$ (early stopping 1000)


Librosa & timbre:



All 41 variables:





Top 10 variables:



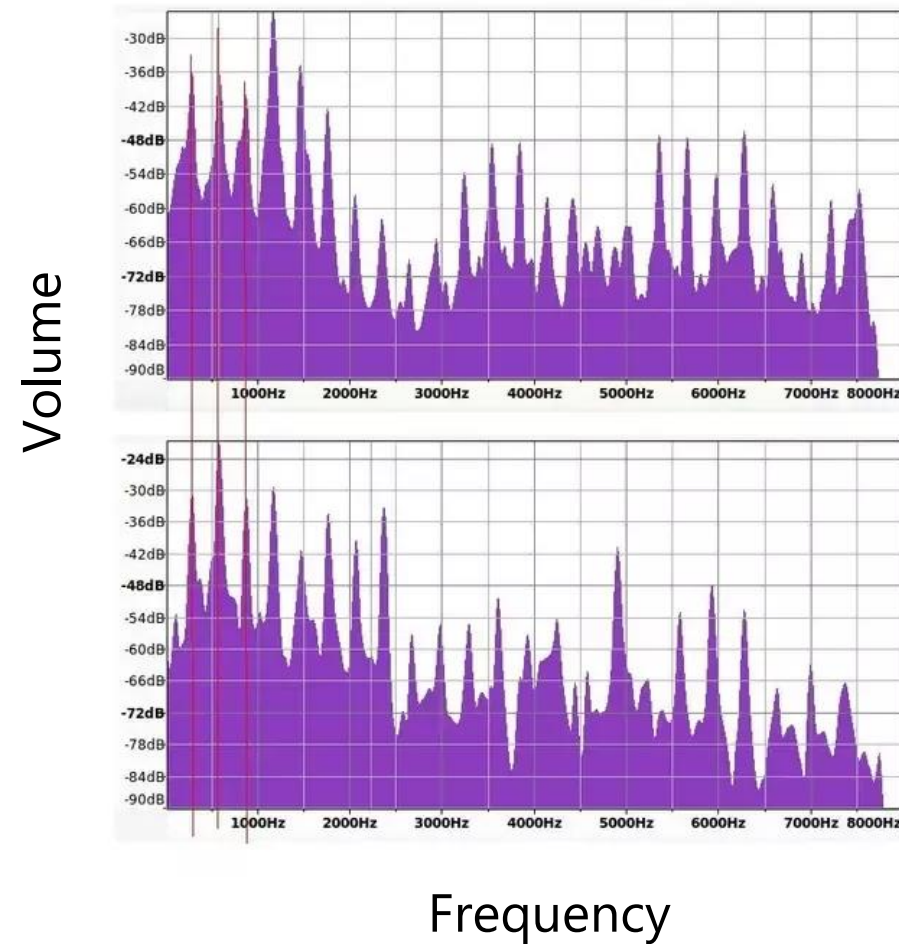
Modest improvement!

- Rythmic features improves model somewhat
- Surprising for me that "Early" & "Late" variables aren't more important
- Many variables, little data...
 - Get more data
 - Dimensionality reduction
 - Recursive feature elimination
- More advanced hyper parameter optimization
- Extract chords and their progressions from Harmonic component


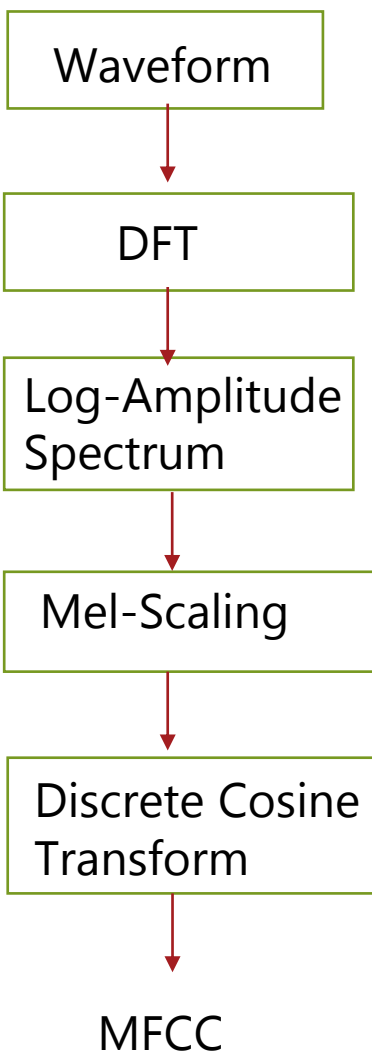


Different approach: Neural network with MFCC's


If you play the same note on a guitar and a piano with the same amplitude, what makes them sound different is **timbre**




Finding features




Example MFCCs




Blues



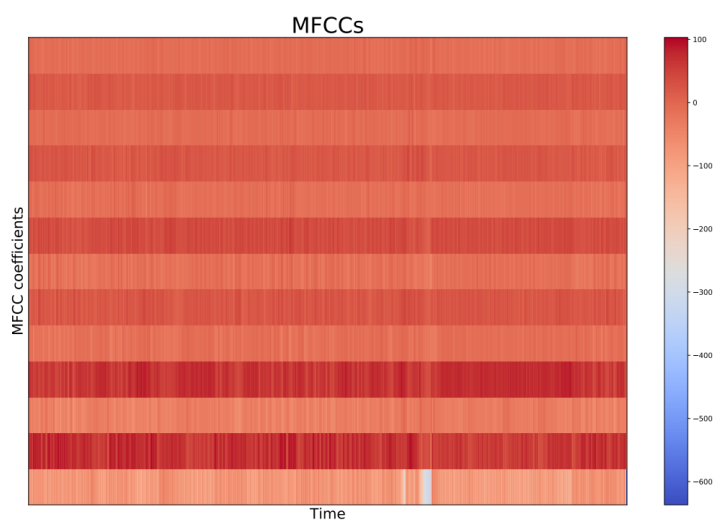
Classical



Country



Rock



Metal


Conventional Neural Network

- Used Keras.Sequential to build CNN
 - 3 x Convolution layer, max pooling, batch normalization
 - 1 Dense layer with 64 neurons
 - Dropout layer
 - Dense layer with 10 neurons
- 30 epochs used
- Takes roughly 30 minutes to train
- Optimized using Keras.Tuner "RandomSearch"
 - number of convolution layers
 - number of filters in each layer

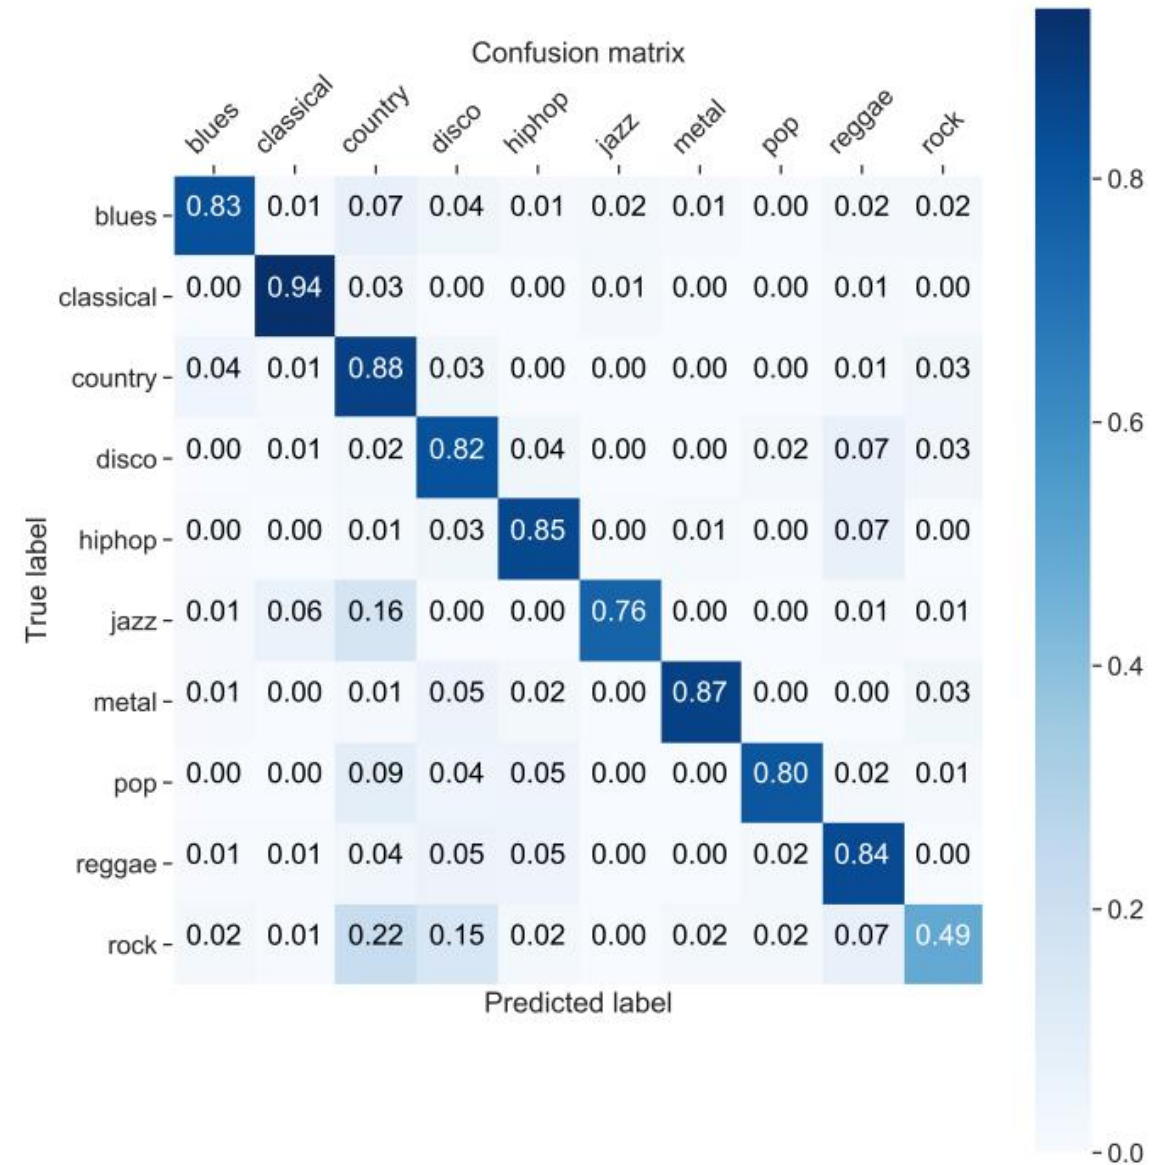
Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 128, 11, 224)	2240
max_pooling2d_6 (MaxPooling2D)	(None, 64, 6, 224)	0
batch_normalization_6 (Batch Normalization)	(None, 64, 6, 224)	896
conv2d_8 (Conv2D)	(None, 62, 4, 224)	451808
max_pooling2d_7 (MaxPooling2D)	(None, 31, 2, 224)	0
batch_normalization_7 (Batch Normalization)	(None, 31, 2, 224)	896
conv2d_9 (Conv2D)	(None, 30, 1, 32)	28704
max_pooling2d_8 (MaxPooling2D)	(None, 15, 1, 32)	0
batch_normalization_8 (Batch Normalization)	(None, 15, 1, 32)	128
flatten_2 (Flatten)	(None, 480)	0
dense_4 (Dense)	(None, 64)	30784
dropout_2 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 10)	650
=====		
Total params: 516,106		
Trainable params: 515,146		
Non-trainable params: 960		

Results



- Scores 80.8% on Test data
- Classical and metal are easiest to predict
- Rock is hardest, mostly mislabeled as country



Outlook


- CNN is a good model for classifying music
- I don't think it will get much better, since humans can't always classify genres perfectly
- Would like to keep optimizing more hyperparameters in the CNN
- Try to extract instruments in all the tracks



Appendix

Feature importance of initial approach

The variables all contribute to the classification. Surprisingly, one of the most important variables is the mean of the fourier transform, done by hand.



Optimizing tree-based algorithm

We use LightGBM for the initial attempts at efficient classification.

The optimal parameters we found to be:

`n_epochs = 500`

`n_leafs = 10`


`max_depth = 10`

This yields a logloss of 0.3 for the multi-classification algorithm, using LightGBM cross validation.

kNearestneighbor

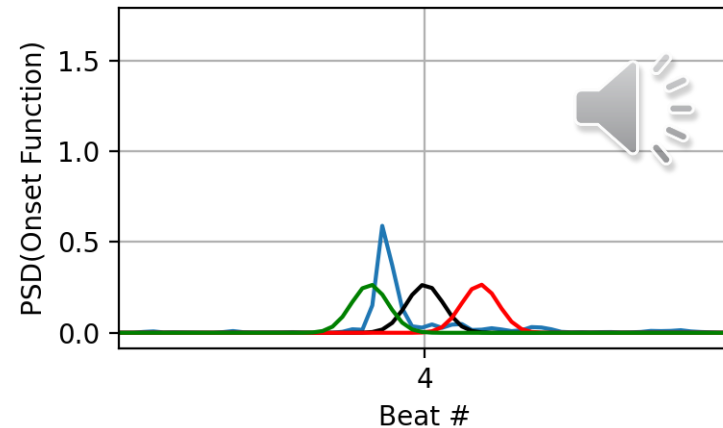
The kNearestneighbor is very straightforward and easy to use. This is the approach that many have used on the data set.

We use `n_neighbors = 12`, based on simple optimization.




Extracted features, training


Early/Late at beat 4 for disco4




Early/Late at beat 2 for blues12



Loss in grid search



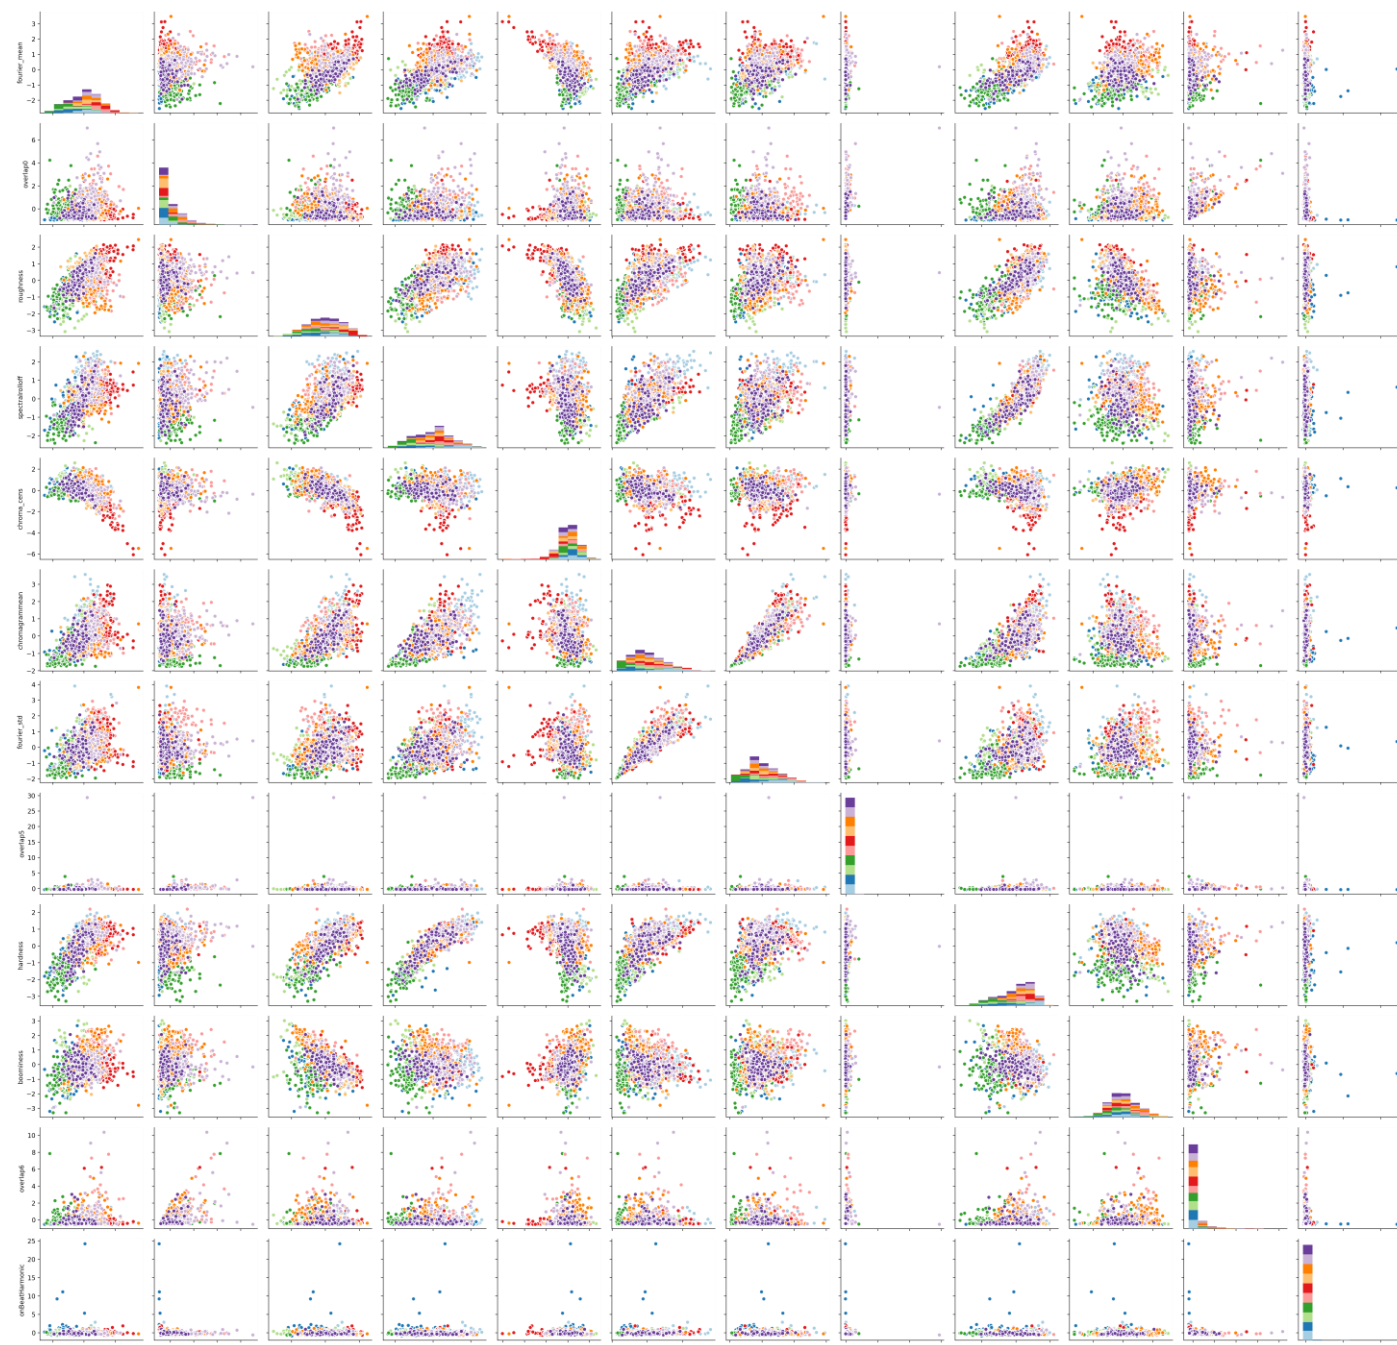
Validation loss best hyper params



Top12 Pairplot

blues
 reggae
 hip-hop
 disco
 rock
 pop
 classical
 country
 jazz
 metal

Fourier_mean
 Overlap0
 Roughness
 Spectralrolloff
 Chroma_cens
 Chromagrammean
 Fourier_std
 Overlap5
 Hardness
 boominess
 Overlap6
 onBeatharmonic



Preprocessing - NN

- Started with 100 samples per genre, then divided each by 5 to have more data
- Extracted MFCC's using the Librosa package
- Parameters: n_fft=2048 (window for fft in num. of samples), hop_length=512 (in num. Of samples)

```
# extract mfcc
mfcc = librosa.feature.mfcc(signal[start:finish], sample_rate, n_mfcc=num_mfcc,
mfcc = mfcc.T
```

- Saved all MFCC's, labels, and mapping to a .json file

```
# save MFCCs to json file
with open(json_path, "w") as fp:
    json.dump(data, fp, indent=4)
```


Tuning Neural Network

```
from kerastuner import RandomSearch
from kerastuner.engine.hyperparameters import HyperParameters
```

- Varied number of filters from 32 to 256 in steps of 32 for 2 Convolutional layers

```
tuner=RandomSearch(build_model,
                   objective='val_accuracy',
                   max_trials=3,
                   directory=LOG_DIR,project_name="MusicClass")
```

```
tuner.search(x=X_train,
            y=y_train,
            epochs=8,
            batch_size=64,
            validation_data=(X_test,y_test))
```

- Tried with 1-3 Convolutional layers, also with 2-3 Dense layers

Neural network details

Activation= ReLu

Max Pool window = 2 x (3,3), then (2,2)

Strides = (2,2)

Padding=same

Optimizer= Adam

Learning rate=0.0001

Loss = sparse_categorical_crossentropy

Metrics=accuracy

```
# build network topology
model = keras.Sequential()

# 1st conv layer
model.add(keras.layers.Conv2D(224, (3, 3), activation='relu', input_shape=input_shape))
model.add(keras.layers.MaxPooling2D((3, 3), strides=(2, 2), padding='same'))
model.add(keras.layers.BatchNormalization())

# 2nd conv layer
model.add(keras.layers.Conv2D(224, (3, 3), activation='relu'))
model.add(keras.layers.MaxPooling2D((3, 3), strides=(2, 2), padding='same'))
model.add(keras.layers.BatchNormalization())

# 3rd conv layer
model.add(keras.layers.Conv2D(32, (2, 2), activation='relu'))
model.add(keras.layers.MaxPooling2D((2, 2), strides=(2, 2), padding='same'))
model.add(keras.layers.BatchNormalization())

# flatten output and feed it into dense layer
model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(64, activation='relu'))
model.add(keras.layers.Dropout(0.3))
#model.add(keras.layers.Dense(32, activation='relu'))
#model.add(keras.layers.Dropout(0.1))
# output layer
model.add(keras.layers.Dense(10, activation='softmax'))
```